

Training dynamics: L_{xt} vs. converted L_{KL} per-ti share

