

Why late timesteps dominate the XOPD per-step loss

A flow-matching account of the `dk` blow-up, artifact vs. signal, and the `xopd_dk_space` reweighting

Flow Factory

Abstract

Theory companion to the empirical observation `docs/xopd/progress/2026-07-12-per-timestep-dk-dominance` (last quarter of steps $\approx 87\%$ of $\sum_k d_k$, final step $\approx 46\%$ at `mean_tail`); the implemented remedy is `train.xopd_dk_space="x0"` (clean-latent loss, §5). Relevant code: `compute_per_step_kl` in `src/flow_factory/trainers/xopd/common.py`, called by `_precompute_d_per_timestep` / `_optimize_train_pass` in `src/flow_factory/trainers/xopd/trainer.py`.

This note defines three implemented per-step MSE knobs (L_{xt} , L_v , L_{x_0}) plus the Gaussian transition KL L_{KL} (DanceOPD theory, §1.1), their exact conversion on the FLUX.2-klein 28-step grid, and why the default L_{xt} (config "`xt`") dominates late timesteps. Companion data: §3–§4.

Contents

1	Three per-step MSE losses and their exact conversion	1
1.1	Fourth form: Gaussian transition KL (DanceOPD / DiffusionOPD theory) . . .	2
2	The key identity (legacy shorthand)	3
3	Concrete schedule: official FLUX.2-klein @ 28 steps	3
4	Empirical distribution: mixed run, converted from measured L_{xt}	4
4.1	Training dynamics: does L_{xt} late dominance change over epochs?	5
5	Implementation: the <code>xopd_dk_space</code> knob	6
6	Two causes of the dominance: separate the artifact from the signal	7
6.1	Cause A — Δt_k^2 amplification in L_{xt} (a parameterization artifact)	7
6.2	Cause B — genuine difficulty near the data manifold (a real signal)	8
6.3	The U-shape	8
7	XOPD-specific reading (large teacher)	8
8	Unified view: every proposal is a choice of per-step weight w_k	9
8.1	Evaluation	9
9	Disambiguating experiment (artifact vs. signal)	9

1 Three per-step MSE losses and their exact conversion

XOPD compares student and teacher on the *same* on-policy rollout state x_t at step index k ; only their predictions differ. Under rectified-flow interpolation ($x_t = (1 - \sigma)z_0 + \sigma\varepsilon$, $\sigma = t/1000$)

and ODE Euler stepping ($\mu = x_t + v \Delta t$, $\Delta t = \sigma_{k+1}^{\text{sched}} - \sigma_k^{\text{sched}} < 0$), there are three equivalent MSE targets (config `xopd_dk_space`, default `"xt"`):

knob	loss	formula	note
<code>xt</code>	L_{xt}	$\ \mu_s - \mu_t\ ^2$	transition mean / next latent (default)
<code>v</code>	L_v	$\ v_s - v_t\ ^2$	raw velocity MSE; ODE-only
<code>x0</code>	L_{x_0}	$\ z_{0,s} - z_{0,t}\ ^2$	clean-latent MSE; ODE-only

Ratio. $\text{MSE}(v) : \text{MSE}(xt) : \text{MSE}(x_0) = 1 : \Delta t^2 : \sigma^2$.

Here μ is the one-step transition mean ($x_{t+\Delta t}$ under ODE), v the model velocity (`noise_pred`), and $z_0 = x_t - \sigma v$ the analytically recovered clean latent. Because x_t cancels:

$$\mu_s - \mu_t = \Delta t (v_s - v_t), \quad (1)$$

$$z_{0,s} - z_{0,t} = -\sigma (v_s - v_t), \quad (2)$$

$$\Rightarrow L_{xt} = \Delta t^2 L_v, \quad (3)$$

$$L_{x_0} = \sigma^2 L_v, \quad (4)$$

$$L_{x_0} = \left(\frac{\sigma}{\Delta t}\right)^2 L_{xt}, \quad (5)$$

$$L_v = \left(\frac{\Delta t}{\sigma}\right)^2 L_{x_0}. \quad (6)$$

All four identities are *exact* on the discrete ODE grid (no approximation). The conversion factors depend *only* on $(\sigma_k, \Delta t_k)$ at step k ; they do not require equal errors across steps — they convert *the same* underlying $(v_s - v_t)$ disagreement between spaces.

Conversion matrix (multiply to go from column loss to row loss).

	$\rightarrow L_{xt}$	$\rightarrow L_v$	$\rightarrow L_{x_0}$
from L_{xt}	1	Δt^{-2}	$(\sigma/\Delta t)^2$
from L_v	Δt^2	1	σ^2
from L_{x_0}	$(\Delta t/\sigma)^2$	σ^{-2}	1

What each loss implicitly weights (if you sum $\sum_k L_{\{\cdot\},k}$ with coefficient 1). Fix the *underlying* velocity disagreement $L_{v,k}$ at step k :

- L_{xt} (**default**): step k contributes $\Delta t_k^2 L_{v,k}$ — late steps have larger $|\Delta t_k|$ and dominate (§3, col. “% L_{xt} ”).
- L_v (`xopd_dk_space="v"`): step k contributes $L_{v,k}$ directly — the native FM velocity target; under the fixed 28-step grid the per-step Δt_k vary only modestly so this is close to uniform (§4).
- L_{x_0} (`xopd_dk_space="x0"`): step k contributes $\sigma_k^2 L_{v,k}$ — early high-noise steps ($\sigma \approx 1$) dominate (col. “pct L_{x_0} ”).

Alternatively, fix $L_{x_0,k}$ (equal clean-latent disagreement): $L_{xt,k} = (\Delta t_k/\sigma_k)^2 L_{x_0,k}$, which is the “ $(\Delta t/\sigma)^2$ blow-up” narrative (§6).

1.1 Fourth form: Gaussian transition KL (DanceOPD / DiffusionOPD theory)

The three knobs above are *unweighted* per-step MSE targets. A separate, widely used objective is the **same-variance Gaussian transition KL** between teacher and student one-step distributions. Under a local model $p(\cdot | x_t) = \mathcal{N}(\mu, \sigma_t^2 I)$ with shared covariance σ_t^2 (DanceOPD Appendix 7.1; DiffusionOPD pathwise term), and $\mu = x_t + v \Delta t$:

$$L_{\text{KL},k} = \text{KL}(p_m(\cdot | x_t) \| p_\theta(\cdot | x_t)) = \frac{\Delta t_k^2}{2\sigma_k^2} \|v_s - v_t\|^2 = \frac{L_{xt,k}}{2\sigma_k^2}. \quad (7)$$

This is **not** equal to any of $\{L_{xt}, L_v, L_{x_0}\}$ with unit coefficient; it is a *fourth* reweighting of the same underlying velocity gap. Exact conversions (same $(v_s - v_t)$ at step k):

	$\rightarrow L_{xt}$	$\rightarrow L_v$	$\rightarrow L_{x_0}$	$\rightarrow L_{\text{KL}}$
from L_{KL}	$2\sigma^2$	$2\sigma^2/\Delta t^2$	$2\sigma^4/\Delta t^2$	1
from L_{xt}	1	Δt^{-2}	$(\sigma/\Delta t)^2$	$(2\sigma^2)^{-1}$
from L_v	Δt^2	1	σ^2	$\Delta t^2/(2\sigma^2)$
from L_{x_0}	$(\Delta t/\sigma)^2$	σ^{-2}	1	$\Delta t^2/(2\sigma^4)$

Implicit weight on L_v . Summing $\sum_k L_{\text{KL},k}$ with coefficient 1 applies $w_k^{\text{KL}} = \Delta t_k^2/(2\sigma_k^2)$ to the underlying $L_{v,k}$ — *between* the L_{xt} weight (Δt_k^2) and the L_{x_0} weight (σ_k^2): early high-noise steps are down-weighted by $1/\sigma_k^2$ relative to L_{xt} , but late steps with small σ_k still *amplify* $1/\sigma_k^2$ enough that L_{KL} is **more** late-heavy than L_{xt} under flat L_v (Figure 1: **77%** vs. **41%** on **ti=27**).

DanceOPD vs. our `normalize_dk`.

- **DanceOPD training default** is *plain* $L_v = \|v_s - v_t\|^2$ at a single semantic-side low- t query — (7) is the *theoretical* KL identity; their ablations report plain MSE beats KL- $\bar{\sigma}^2$ / Min-SNR reweightings (`docs/mof/danceopd-comparison.md`).
- **Our `normalize_dk=true`** (only in `space="xt"`) divides transition mean MSE by $2\bar{\sigma}_k^2 = 2\text{std_dev}_t^2 |\Delta t_k|$ (SDE Appendix B $\bar{\sigma}$), *not* by $2\sigma_k^2$ with FM fraction $\sigma_k = t_k/1000$. Under ODE ($\text{std_dev}_t \approx 0$) it falls back to plain L_{xt} . So (7) is **conceptually related** to normalized transition KL but **not implemented** as a separate knob; it differs from L_{KL} unless one identifies $\bar{\sigma}$ with σ_k and uses SDE dynamics.
- **Ratio to L_{xt} :** $L_{\text{KL},k}/L_{xt,k} = 1/(2\sigma_k^2)$. On the FLUX.2-klein grid this *increases* late-step share vs. plain L_{xt} (77% vs. 41% on **ti=27** under flat L_v), because $\sigma_k \rightarrow 0$ faster than Δt_k^2 shrinks toward the end.

2 The key identity (legacy shorthand)

With `normalize_dk=false`, the logged `train/dk/{ti}` is L_{xt} (not L_v). Combining (3)–(6):

$$L_{xt} = \Delta t^2 L_v = \left(\frac{\Delta t}{\sigma}\right)^2 L_{x_0}.$$

So the default loss is clean-latent disagreement multiplied by the implicit per-step weight $(\Delta t/\sigma)^2$ (§1). The x -space knob sets L_{x_0} directly and removes that factor (§5).

3 Concrete schedule: official FLUX.2-klein @ 28 steps

We reproduce the *exact* denoising grid used in XOPD training (`num_inference_steps=28`, 512×512) by loading the default scheduler from the official HuggingFace repo and calling the same `retrieve_timesteps` path as `Flux2KleinPipeline` (shifted flow-matching sigmas with `image_seq_len=1024`, $\mu = 1.859$):

```
FlowMatchEulerDiscreteScheduler.from_pretrained(
    "black-forest-labs/FLUX.2-klein-base-4B", subfolder="scheduler")
# sigmas = linspace(1, 1/28, 28); mu = compute_empirical_mu(1024, 28)
```

Repro script: `.scratch/flux2_klein_scheduler_timesteps.py` (uv venv /Users/bowenping/.venvs/flux2-diffusers==0.39.0).

Notation at step index `ti`. `ti=0` is the noisiest state; `ti=27` is the last denoising step before $\sigma = 0$. The scheduler returns a *timestep* $t_k \in [0, 1000]$ and a *shifted flow sigma* $\sigma_k^{\text{sched}} \in [0, 1]$. XOPD’s clean-latent recovery uses the repo’s FM fraction $\sigma_k = t_k/1000$ (`_noise_fraction`, `flow_match_sigma`). The ODE step size is $\Delta t_k = \sigma_{k+1}^{\text{sched}} - \sigma_k^{\text{sched}} < 0$ (from `scheduler.step`). On this grid $\sigma_k \approx \sigma_k^{\text{sched}}$ to machine precision.

Per-step conversion factors on the 28-step grid. Table 1 lists representative steps; Table 2 (auto-generated, `docs/xopd/flux2_klein_28step_loss_scales.tex`) gives all 28 rows. Repro: `.scratch/flux2_klein_scheduler_timesteps.py`.

<code>ti</code>	t	σ	Δt	$\frac{L_{xt}}{L_v} = \Delta t^2$	$\frac{L_{x_0}}{L_v} = \sigma^2$	$\frac{L_{x_0}}{L_{xt}}$	pct L_{xt}	pct L_{x_0}
0	1000.0	1.0000	-0.00574	3.3×10^{-5}	1.000	3.0×10^4	0.04%	5.28%
7	950.6	0.9506	-0.00929	8.6×10^{-5}	0.904	1.0×10^4	0.10%	4.77%
13	881.0	0.8810	-0.01584	2.5×10^{-4}	0.776	3.1×10^3	0.28%	4.10%
21	681.5	0.6815	-0.04505	2.0×10^{-3}	0.464	2.3×10^2	2.27%	2.45%
24	516.8	0.5168	-0.08175	6.7×10^{-3}	0.267	40.0	7.47%	1.41%
25	435.1	0.4351	-0.10456	1.1×10^{-2}	0.189	17.3	12.21%	1.00%
26	330.5	0.3305	-0.13847	1.9×10^{-2}	0.109	5.70	21.42%	0.58%
27	192.1	0.1921	-0.19206	3.7×10^{-2}	0.037	1.00	41.21%	0.19%

Table 1: FLUX.2-klein-base-4B, `num_inference_steps=28`, 512px. Columns 5–7: exact factors from (3)–(5). “% L_{xt} ” (“% L_{x_0} ”): share of $\sum_k L_{xt,k}$ (resp. $\sum_k L_{x_0,k}$) if $L_{v,k} \equiv \text{const}$ across steps.

Reading the numbers.

- **At `ti=27`, all three losses coincide in scale.** $|\Delta t_{27}| \approx \sigma_{27}$, so $\Delta t_{27}^2 = \sigma_{27}^2 \approx 0.037$ and $L_{xt} = L_{x_0} = L_v$ for the *same* velocity gap.
- **Default L_{xt} dominance $\approx$ flat L_v .** If $\|v_s - v_t\|^2$ were identical at every step, $\sum_k L_{xt,k}$ puts 41% on `ti=27` (col. “% L_{xt} ”). The mixed XOPD run observes **46%** on `ti=27` at `mean_tail` — the logged curve matches *roughly constant velocity error* amplified only by Δt_k^2 , not the stronger $(\Delta t_k/\sigma_k)^2$ factor (which would give **77%** if $L_{x_0,k}$ were constant).
- **L_{x_0} up-weights early steps.** With flat L_v , high-noise steps each take $\sim 5\%$ of $\sum L_{x_0}$ while `ti=27` takes only **0.19%**. Setting `xopd_dk_space="x0"` switches to this allocation.
- **Cross-step ratios (same L_v).** $L_{xt,27}/L_{xt,13} = (\Delta t_{27}/\Delta t_{13})^2 \approx 148\times$; $L_{x_0,0}/L_{x_0,27} = \sigma_0^2/\sigma_{27}^2 \approx 27\times$.
- **“Late step” $\neq \sigma \rightarrow 0$.** The last rollout step has $t = 192$ ($\sigma = 0.19$), not $t \approx 0$. Dominance comes from the discrete grid’s large $|\Delta t_k|$ at the end.

ti	t	σ	Δt	$\frac{L_{xt}}{L_v}$	$\frac{L_{x_0}}{L_v}$	$\frac{L_{x_0}}{L_{xt}}$	pct L_{xt}	pct L_{x_0}
0	1000.0	1.0000	-0.00574	3.3e-05	1.000	30380.38	0.04%	5.28%
1	994.3	0.9943	-0.00611	3.7e-05	9.89e-01	26518.80	0.04%	5.22%
2	988.2	0.9882	-0.00651	4.2e-05	9.76e-01	23039.00	0.05%	5.15%
3	981.6	0.9816	-0.00696	4.8e-05	9.64e-01	19912.21	0.05%	5.09%
4	974.7	0.9747	-0.00745	5.6e-05	9.50e-01	17114.06	0.06%	5.01%
5	967.2	0.9672	-0.00800	6.4e-05	9.36e-01	14621.44	0.07%	4.94%
6	959.2	0.9592	-0.00861	7.4e-05	9.20e-01	12410.80	0.08%	4.86%
7	950.6	0.9506	-0.00929	8.6e-05	9.04e-01	10460.53	0.10%	4.77%
8	941.3	0.9413	-0.01006	1.0e-04	8.86e-01	8748.98	0.11%	4.68%
9	931.3	0.9313	-0.01093	1.2e-04	8.67e-01	7256.11	0.13%	4.58%
10	920.3	0.9203	-0.01192	1.4e-04	8.47e-01	5962.25	0.16%	4.47%
11	908.4	0.9084	-0.01305	1.7e-04	8.25e-01	4849.23	0.19%	4.35%
12	895.4	0.8954	-0.01434	2.1e-04	8.02e-01	3899.25	0.23%	4.23%
13	881.0	0.8810	-0.01584	2.5e-04	7.76e-01	3095.63	0.28%	4.10%
14	865.2	0.8652	-0.01758	3.1e-04	7.49e-01	2422.59	0.35%	3.95%
15	847.6	0.8476	-0.01963	3.9e-04	7.18e-01	1865.26	0.43%	3.79%
16	828.0	0.8280	-0.02205	4.9e-04	6.86e-01	1409.55	0.54%	3.62%
17	805.9	0.8059	-0.02496	6.2e-04	6.50e-01	1042.44	0.70%	3.43%
18	781.0	0.7810	-0.02849	8.1e-04	6.10e-01	751.66	0.91%	3.22%
19	752.5	0.7525	-0.03281	1.08e-03	5.66e-01	525.93	1.20%	2.99%
20	719.7	0.7197	-0.03821	1.46e-03	5.18e-01	354.83	1.63%	2.73%
21	681.5	0.6815	-0.04505	2.03e-03	4.64e-01	228.85	2.27%	2.45%
22	636.4	0.6364	-0.05391	2.91e-03	4.05e-01	139.37	3.25%	2.14%
23	582.5	0.5825	-0.06567	4.31e-03	3.39e-01	78.683	4.82%	1.79%
24	516.8	0.5168	-0.08175	6.68e-03	2.67e-01	39.970	7.47%	1.41%
25	435.1	0.4351	-0.10456	1.09e-02	1.89e-01	17.315	12.21%	1.00%
26	330.5	0.3305	-0.13847	1.92e-02	1.09e-01	5.698	21.42%	0.58%
27	192.1	0.1921	-0.19206	3.69e-02	3.69e-02	1.000	41.21%	0.19%

Table 2: Three per-step MSE losses on the FLUX.2-klein 28-step grid (512×512). Columns 5–7 are exact conversion factors at step k (on-policy, shared x_t). Last two columns: normalized share of $\sum_k L_{xt}$ (resp. $\sum_k L_{x_0}$) if $\|v_s - v_t\|^2$ were identical at every step. Generated from `black-forest-labs/FLUX.2-klein-base-4B` default scheduler.

Implicit per-step shares (flat L_v). Figure 1 plots the normalized share of $\sum_k L_{\{ \cdot \}, k}$ at each `ti` when the underlying velocity gap is identical at every step ($L_{v,k} \equiv \text{const}$). Blue/red/green: the three implemented MSE knobs (L_{xt} , L_v , L_{x_0}); orange: the Gaussian transition KL $L_{\text{KL}} = \Delta t^2 L_v / (2\sigma^2)$ (§1.1). Its weight $\Delta t_k^2 / (2\sigma_k^2)$ is *more* late-heavy than L_{xt} alone under flat L_v .

4 Empirical distribution: mixed run, converted from measured

L_{xt}

We pull `train/dk/{ti}` from two W&B runs in project `Flow-Factory-XOPD`, both with `xopd_dk_space=xt` and `normalize_dk=false`: `4cnkluwk` (mixed geneval+OCR 1:1) and `51ueyo7w` (geneval.enhanced specialist). Under ODE we convert each step to the *hypothetical* L_v and L_{x_0} that would have been logged with the same underlying velocity gap:

$$L_{v,k} = \frac{L_{xt,k}}{\Delta t_k^2}, \quad L_{x_0,k} = \left(\frac{\sigma_k}{\Delta t_k} \right)^2 L_{xt,k}, \quad L_{\text{KL},k} = \frac{L_{xt,k}}{2\sigma_k^2}.$$

Windows. *initial* = mean of the first 4 logged optimizer steps ($\approx$ first 4 epochs); *mean_tail* = mean of the last 30% of logged steps. Figures are generated by `.scratch/plot_per_timestep_loss_dominance`. (repro: pull W&B, then plot).

Table 3 summarizes the mixed run `mean_tail` aggregates; values match Figure 2 (top-right panel).

Reading.




Figure 1: Theoretical per-ti shares under flat- L_v . Four curves: L_{xt} , L_v , L_{x_0} , and DanceOPD/DiffusionOPD transition KL $L_{KL} = \Delta t^2 / (2\sigma^2) \cdot L_v$. L_{KL} (orange) is **more** late-concentrated than L_{xt} (77% vs. 41% on $\text{ti}=27$) because $1/\sigma_k^2$ diverges as $\sigma_k \rightarrow 0$.

aggregate	L_{xt} (meas.)	L_v (conv.)	L_{x_0} (conv.)
top-1 ($\text{ti}=27$)	45.6%	0.5%	0.02%
top-7 ($\text{ti}=21\text{--}27$, last 1/4)	87.3%	3.4%	0.9%
high-noise half ($\text{ti}=0\text{--}13$)	6.9%	85.9%	92.1%
low-noise half ($\text{ti}=14\text{--}27$)	93.1%	14.1%	7.9%

Table 3: Share of $\sum_k L_{\{\cdot\},k}$ (`mean_tail`, mixed run `4cnkluwk`, 250 logged steps) in each loss space. L_{xt} is measured; L_v and L_{x_0} are ODE-exact conversions. Geneval_enh specialist initial window: $\text{ti}=27$ **42%**, $\text{ti}=21\text{--}27$ **79%** (`51ueyo7w`).

- **Late dominance is specific to L_{xt} .** Mixed `mean_tail` puts **46%** on $\text{ti}=27$ and **87%** on the last quarter for L_{xt} ; converted L_v and L_{x_0} put **86%/92%** on the high-noise half ($\text{ti}=0\text{--}13$). The logged curve is dominated by Δt_k^2 , while the underlying velocity gap is larger at high noise (Figure 2).
- **Initial $\approx$ theoretical flat- L_v prediction.** Geneval_enh at epoch 0–3 already shows **42%** on $\text{ti}=27$ for L_{xt} , close to the **41%** from Figure 1 — the dominance is largely a metric artifact, not something training must “learn” first.
- **L_v is nearly flat in absolute scale** after conversion (only **0.5%** of the sum on $\text{ti}=27$), consistent with Cause A (§6.1).
- **Ablation implication.** `xopd_dk_space="v"` removes the Δt_k^2 rescaling and should redistribute gradient toward high-noise steps; `"x0"` applies σ_k^2 and amplifies that further.

4.1 Training dynamics: does L_{xt} late dominance change over epochs?

Figure 4 tracks the L_{xt} and converted L_{KL} shares at $\text{ti}=27$, the last quarter ($\text{ti}=21\text{--}27$), and the high-noise half ($\text{ti}=0\text{--}13$, L_{xt} only) across logged optimizer steps (one step $\approx$ one epoch under XOPD). Both runs start with high late-step L_{xt} share ($\sim 40\%$ on $\text{ti}=27$) while L_{KL} is




Figure 2: Empirical per-ti shares in each loss space. Top row: mixed run (4cnkluwk); bottom row: geneval_enh specialist (51ueyo7w). Left: initial window; right: `mean_tail`. Orange: converted $L_{KL} = L_{xt}/(2\sigma_k^2)$ (§1.1). Late dominance is strongest in L_{KL} , absent in L_v/L_{x_0} .

higher still ($\sim 70\text{--}80\%$); both remain stable as training progresses. Bottom panel: gap $L_{KL} - L_{xt}$ at `ti=27` (percentage points), showing the extra late weight from $1/(2\sigma_k^2)$.

5 Implementation: the `xopd_dk_space` knob

Knob. `train.xopd_dk_space` $\in \{v, xt, x_0\}$ (default `xt`). See `compute_per_step_kl` in `xopd/common.py`; `v` and `x_0` require ODE (`XOPDTrainer.__init__` guard).

knob	per-step loss	relation to L_v
<code>xt</code>	$L_{xt} = \ \mu_s - \mu_t\ ^2$	$L_{xt} = \Delta t^2 L_v$
<code>v</code>	$L_v = \ v_s - v_t\ ^2$	identity
<code>x_0</code>	$L_{x_0} = \ z_{0,s} - z_{0,t}\ ^2$	$L_{x_0} = \sigma^2 L_v$

Since x_t cancels, $L_{x_0} = (\sigma/\Delta t)^2 L_{xt}$ ((5)). `normalize_dk` applies only to `xt`; `v`/`x_0` ignore it.

No extra cost. `x_0` recovers z_0 analytically from μ ; `v` uses $(\mu_s - \mu_t)/\Delta t$. **Logging.** `train/dk/{ti}` values are in the selected space; compare *shapes* across spaces, not raw magnitudes (§4).




Figure 3: Top row: measured $L_{xt,k}$ (log- y); bottom row: converted $L_{KL,k} = L_{xt,k}/(2\sigma_k^2)$. Both runs show the L_{xt} U-shape; L_{KL} is even more concentrated on late $\mathbf{ti}$ because of the $1/\sigma_k^2$ factor.

6 Two causes of the dominance: separate the artifact from the signal

6.1 Cause A — Δt_k^2 amplification in L_{xt} (a parameterization artifact)

The default loss is $L_{xt} = \Delta t^2 L_v$ ((3)), not velocity MSE. Summing $\sum_k L_{xt,k}$ with unit weight therefore applies a per-step factor Δt_k^2 to the underlying velocity gap. Late rollout steps have larger $|\Delta t_k|$ (Table 1), so L_{xt} **over-represents late steps even if L_v were flat**.

Quantitative check on the FLUX.2-klein 28-step grid (§3):

- If $L_{v,k} \equiv \text{const}$: $\sum_k L_{xt,k}$ puts **41%** on $\mathbf{ti}=27$ (col. “% L_{xt} ”). The mixed XOPD run observes **46%** on $\mathbf{ti}=27$ for logged L_{xt} — the curve matches *roughly constant velocity error* $\times \Delta t_k^2$, not a separate “mystery” effect.
- If instead $L_{x_0,k} \equiv \text{const}$: $L_{xt,k} = (\Delta t_k/\sigma_k)^2 L_{x_0,k}$ would put **77%** on $\mathbf{ti}=27$. The observed **50 $\times$** ratio $L_{\mu,27}/L_{\mu,13}$ is *between* the flat- L_v prediction ($\sim 148\times$ from Δt ratio alone, before L_v shrinkage) and the flat- L_{x_0} prediction ($\sim 3100\times$), because both the metric *and* the actual L_v vary with k .
- Switching to L_{x_0} (`xopd_dk_space="x0"`) replaces the Δt_k^2 weighting with σ_k^2 (col. “% L_{x_0} ”), up-weighting high-noise structure steps.

This parallels diffusion training, where ε -prediction MSE needs SNR reweighting to avoid low-noise dominance; here the analogue of “native” FM regression is L_v , not L_{xt} .

6.2 Cause B — genuine difficulty near the data manifold (a real signal)

Near clean, the posterior $\mathbb{E}[z_0 \mid x_t]$ sharpens, the velocity field’s spatial Jacobian (Lipschitz constant) grows $\sim 1/\sigma$, and the target encodes sharp, high-frequency, multimodal detail. Under `normalize_dk=false` the latent scale is also largest at clean. This part of the late-step cost is **real** and tied to final-image quality.

6.3 The U-shape

The observation doc’s minimum at $ti \approx 12\text{--}13$ and mild high-noise bump at $ti = 0$ follow from $L_{xt} = \Delta t^2 L_v$ with L_v largest at high noise: early steps have small Δt^2 but large L_v ; late steps have large Δt^2 but shrinking L_v ; the minimum is where these cross.

Corollary: `normalize_dk=true` makes it worse, not better

`normalize_dk=true` divides by $2\bar{\sigma}^2 = 2\text{std_dev}_t^2 |\Delta t|$, and $\bar{\sigma} \rightarrow 0$ toward clean, stacking a *second* amplification on top of $1/\sigma^2$. This is why transition-KL objectives (DanceGRPO/DDPO-style) intrinsically favor the low-noise region. **To equalize timesteps you must go the opposite way.**

7 XOPD-specific reading (large teacher)

“XOPD benefits more because the teacher is larger” is correct, but the benefit lives in **Cause B**, not A:

- The 32B→4B teacher’s advantage (fine detail, texture, aesthetics) concentrates in the *low-noise* steps, where on-policy x_t is already near the student’s final sample — the most valuable place for a big teacher to “polish”.
- Structural/semantic correctness (geneval count/position, OCR) is decided at *high-noise* steps.

So “large teacher” does **not** imply “train only late steps”. It implies “the late-step Cause-B component is worth keeping — but do not mistake the Cause-A amplification for that value.”

8 Unified view: every proposal is a choice of per-step weight w_k

All proposals choose how to weight the underlying velocity gap $L_{v,k}$ across steps: $\mathcal{L} = \sum_k w_k L_{v,k}$ (equivalently $\sum_k w'_k L_{\{xt,v,x_0,KL\},k}$).

direction	loss / weight on L_v	effect
default	$L_{xt}, w_k = 1 \Rightarrow w_k L_v = \Delta t_k^2$	late-step dominance (41% on $ti=27$ if L_v flat)
transition KL (theory)	$L_{KL}, w_k = \Delta t_k^2 / (2\sigma_k^2)$	stronger late dominance than L_{xt} (77% on $ti=27$); DanceOPD trains plain L_v instead
1. late-only (DanceOPD)	L_{xt} or L_v on $\sigma_k < \tau$ only	polish; drops structure learning
2a. L_{x_0} (implemented)	$L_{x_0}, w_k = 1 \Rightarrow w_k L_v = \sigma_k^2$	up-weights high-noise steps
2b. L_v (implemented)	<code>xopd_dk_space="v"</code>	native FM velocity MSE
2c. EMA balance	$L_{xt}/\text{EMA}(L_{xt})$	flattens logged scale; may suppress Cause B
3. endpoint-sensitivity	weight $\propto \partial z_0 / \partial v_k$	2a is first-order proxy

8.1 Evaluation

- **Direction 1 (late-only)** trains L_{xt} on a late window only; same Δt_k^2 factor within the window. Drops high-noise steps where L_{x_0} (hence structure) matters most.
- **Direction 2a** (L_{x_0}) replaces Δt_k^2 with σ_k^2 on the shared L_v . **Direction 2b** (L_v) removes the Δt_k^2 factor entirely (§4). **Do not use `normalize_dk=true`** to equalize (§6.3).
- **Direction 3 (sensitivity)** idealizes weighting L_v by endpoint Jacobian; L_{x_0} is a good first-order proxy.

Recommendation

Compare `xopd_dk_space` $\in \{\text{xt}, \text{v}, \text{x0}\}$ on OCR (§9): default `xt` vs. velocity `v` vs. clean-latent `x0`, plus late-only `xt` (direction 1).

9 Disambiguating experiment (artifact vs. signal)

The ablation in `docs/xopd/progress/2026-07-12-ablation-relay.md` (all on OCR, otherwise identical) is exactly the test:

variant	config	reads
full-timestep xt-MSE	<code>flux2_klein.32b.to.4b.l1.ocr.1kep.yaml</code>	default (xt)
late-timestep xt-MSE	<code>flux2_klein.32b.to.4b.l1.ocr.selective.teaching.1kep.yaml</code>	direction 1
full-timestep x0-MSE	<code>flux2_klein.32b.to.4b.l1.ocr.xspace.1kep.yaml</code>	direction 2a
full-timestep v-MSE	<code>flux2_klein.32b.to.4b.l1.ocr.vmse.1kep.yaml</code>	direction 2b

Interpretation:

- If `x0` or `v` **improves** high-noise-half learning and structure metrics *without* degrading clean-end detail $\rightarrow$ the original `xt` dominance was mostly **Cause A (artifact)**; reweighting is a net win.
- If `x0/v` **hurts** clean-end detail with insufficient structure gains $\rightarrow$ the late-step **Cause B (real value)** dominates; consider soft late-window `xt`.

Remark. Logged $\text{train}/\text{d}_k/\{ti\}$ magnitudes differ across $v/\text{xt}/\text{x0}$; compare per-timestep shape and downstream eval (§4), not raw values across spaces.




Figure 4: L_{xt} (run color) vs. converted L_{KL} (orange, same line style per run) per-timestep share vs. epoch. Mixed (4cnkluwk, solid) and geneval_enh (51ueyo7w, dashed). Bottom: Δ ti=27 share ($L_{KL} - L_{xt}$).